

P56 - Lab 5 Performance with Antagonists

Professor Andrew W. Moore (and explicit thanks to Dr Gary Guo)

Lent, 2023/2024

The aim of this lab is

- to extend previous Kernel tracing work by exploring the impact of an antagonistic workload upon the previous lab workloads.

The approach for this lab will be to rerun a number of experiments from Lab 3 – the server.disk experiments of Book Section 7.7 (exercise 3 (writing), exercise 4 (reading for checksum)), and exercise 5 causing server stress with multiple clients; studying this application using the KUtrace system.

1 Operating with an Antagonist

The use of an antagonist can identify specific problem areas; by putting a component under pressure (memory, CPU, disk, network, etc.) we can identify both immediate issues and points with potential for starvation.

For these experiments repeat Lab 3's experiment 3, 4, and 5 again but this time with multiple copies of an antagonist. Our antagonist will be a cpuhog from `~/KUtrace/bookcode/aw_files/book-user-code`

Code for the cpuhog can be (re)compiled using the command:

```
(code in ~/KUtrace/bookcode/aw_files/book-user-code/cpuhog_highissue.cc)
g++ cpuhog_highissue.cc -o cpuhog
```

cpuhog will exit after about 50 minutes.

You may uncover interesting behaviours as you run experiments with varying numbers of simultaneous cpuhog executables.

An example command to run three cpuhog processes with the server.

```
~/KUtrace/bookcode/aw_files/book-user-code/cpuhog & \
~/KUtrace/bookcode/aw_files/book-user-code/cpuhog & \
~/KUtrace/bookcode/aw_files/book-user-code/cpuhog & \
~/KUtrace/bookcode/aw_files/book-user-code/server_disk diskdir
```

Further CPU pressure can be achieved by increasing the number of cpuhog processes and by changing the scheduling priority. This is an example using the `nice(1)` command to increase run priority for each of the three cpuhog processes. You may wish to experiment with other combinations to trigger an interesting (study-worthy) effect.

```
cat diskdir/a* diskdir/c* diskdir/d* > /dev/null
sudo nice --10 ~/KUtrace/bookcode/aw_files/book-user-code/cpuhog & \
sudo nice --10 ~/KUtrace/bookcode/aw_files/book-user-code/cpuhog & \
sudo nice --10 ~/KUtrace/bookcode/aw_files/book-user-code/cpuhog & \
~/KUtrace/bookcode/aw_files/book-user-code/server_disk diskdir
```

Experiment 3 from Lab 3 uses the data created in previous experiments (the data in `diskdir`).

These commands will ensure the experiment starts from a consistent state; first by purging the operating-system cache then by systematically accessing the `a*` and `d*` files prior to the experiment. This has the result of imposing the disk load overhead on the `c*` files only.

For each set of experiments, this will provide better repeatability and also by starting in a known state, your experiments become more consistent too.

```
# on server
# 151-piAAA
# Purge page cache
echo 1 | sudo tee /proc/sys/vm/drop_caches

# Bring files back to RAM
cat diskdir/aaaa diskdir/ddddd > /dev/null
```

refreshing the cache, thereby imposing the disk overhead on all key accesses.

2 Understanding Your Experiments

Remember to start each experiment with a model of your understanding, then run the experiment and you can either validate or improve, develop, and refine your understanding.

By this stage you will have run multiple experiments and followed these observations with exploration of behaviour varying as it is subject to antagonistic behaviour from cpuhog processes. Identify behaviours you have been alerted to in the textbook, three examples can include lock contention, cache misuse, or disk contentions; other issues are also present. These issues and further issues will develop when subjected to cpu contention and cpu exhaustion.

It is precisely such situations that lead to poorer behaviour and that may be identified with tools such as KUtrace. For each experiment look to create (in the antagonistic) or

simply to identify one such performance issues, to document the issues and to clearly demonstrate your understanding of the wealth of tracing information presented.

The following prompts are intended to help you understand your results, while you also provide supporting evidence for your report. However, they are just suggestions — feel free to approach the data differently!

- Discuss the methodology of the measurement tools.
- Explain how the limitations of the methodology are mitigated in this lab.
- Explore the limitations of the experiments conducted in this lab, and explain where the quality of the experiment (e.g., setup, methodology) could have been improved or altered.

Finally, you should always look for odd or surprising results (such as time spent in unexpected ways), and try to explain them. Note that sometimes exceptional results indicate a problem in your setup or scripts.